
DRAFT

CMS Internal Note

The content of this note is intended for CMS internal use and distribution only

Disk Storage Usage and Metrics for Dynamic Data Management

S. Narayanan, M. Goncharov, and C. Paus
Massachusetts Institute of Technology

Abstract

In the CMS experiment, detector and Monte Carlo simulation data are ordered in datasets, which have some common properties and are usually analyzed as a whole. The popularity of such datasets varies substantially, which opens the question how to best distribute the datasets such that they are optimally accessible at the various computing sites. Generally speaking, popular datasets should be replicated at several sites while less popular datasets might just have a single copy in the overall system. Also as the data taking progresses and new Monte Carlo simulation datasets become available, those datasets have to be distributed in the system and outdated datasets have to be removed. The Dynamic Data Management tools automatically manage the replication of datasets in the distributed multi-site computing system with the goal of optimising the system performance. In this note, we describe a metric for the performance of Dynamic Data Management, based on the number of user accesses of a the given datasets.

This box is only visible in draft mode. Please make sure the values below make sense.

PDFAuthor:	S. Narayanan, M. Goncharov, C. Paus
PDFTitle:	Disk Storage Usage and Metrics for Dynamic Data Management
PDFSubject:	CMS
PDFKeywords:	CMS, Dynamic Data Management, popularity

Please also verify that the abstract does not use any user defined symbols

Contents

1	1	Introduction	1
2	2	Ingredients	1
3	2.1	Dataset selection	1
4	2.2	Average N_{replicas}	2
5	2.3	Number of Accesses N_{accesses} and File N_{files} , and Size size	2
6	3	Filling the plot	3
7	3.1	Dataset selection	3
8	3.2	Binning	3

1 Introduction

Dynamic Data Management currently manages a pool of approximately 20 PB across several Tier-2 and Tier-1 sites. The purpose of the management to first order is the creation and deletion of replicas of datasets that are part of the pool. Additional replicas of a dataset are created when the dataset is particular popular while excess replicas are deleted when they are less popular. The creation and deletion of replicas does not add new datasets to the pool nor does it remove datasets entirely from it. Adding new datasets to the pool and completely removing datasets from it is treated separately. Newly created datasets that are relevant to the Physics Group get automatically added to the pool by the Computing Operations group and datasets will be deleted following deletion campaigns that are driven by the Physics Groups in the experiment. This means usually for a given dataset there is at least one copy in the system often referred to as the last copy. If a dataset is declared deprecated, all copies will be deleted.

A good measure of the performance of the algorithms used to create and delete dataset replicas is the number of accesses per replica. If, for a given dataset, this number is very large, then the dataset is not sufficiently replicated. On the other hand, if it is very small for many datasets, then we are maintaining too many replicas of unused datasets. To produce plots of the popularity of datasets in a given time interval we have to carefully determine the list of datasets that we have to consider and determine four attributes for each dataset: number of accesses, size on disk, number of files, and average number of replicas.

The plot provides a measure of how well our computing system is using the given disk space in our data pool. In the following we explain in detail how we determine the list of datasets we consider and how we calculate the four above listed properties for each dataset.

2 Ingredients

2.1 Dataset selection

The CMS experiment creates a large number of datasets of which not all are commonly used by the people doing analysis. From the detector we have RAW and RECO data formats, but since 2011 the much more compressed AOD datasets had been established as the datasets used for analysis. The other data formats are mostly stored on tape and are used for production purposes when the data or Monte Carlo simulation samples get re-reconstructed with better calibrations/alignments and/or new reconstruction algorithms. The majority of data on the Tier-2 centers are therefore of AOD or AODSIM format or since recently the even more reduced MINIAOD and MINIAODSIM format, but there are some other formats. In this study we

evaluate data popularity in a given time interval and distinguish between two scenarios for the selection of datasets to consider.

- include all datasets that are managed by phedex ?? in any of the official data groups and that were at some point during the considered interval at a Tier-2 site, or
- include all datasets that are or were in the past managed by phedex ?? in the 'AnalysisOps' data group (this is the dynamic data management pool) and that were at some point during the considered interval at a Tier-2 site.

We distinguish these two scenarios because the first one very generally shows the usage of our entire storage, while the second scenario zooms in on the dynamically managed data pool, and thus offers itself as a tool to optimize the data distribution for the best system performance.

While the dataset selection appears to be a straight forward the fact that we do not have a full record of the contents of the data storage of the order 50 Tier-2 sites makes it complicated. Obtaining the list of relevant datasets is a non trivial issue. In the popularity database we can easily find all datasets that were used at a particular site in the given time interval, but datasets that were not accessed during this time interval are not recorded and therefore we must use an alternative source to find all datasets.

2.2 Average N_{replicas}

Phedex directly provides us with the current locations of all datasets. However, this information is not directly available for the past. Thus, Phedex transfer and deletion histories are used to infer the timeline of the presence of a dataset on a given site.

The histories are 'sanitized' to remove self-inconsistent entries such as the transfer of a dataset to a site on which it already exists. It is assumed that each site can only contain one copy of a dataset. If there is no Phedex history for a given dataset on a given site, but we know that the dataset is currently on that site, then it is assumed to have existed on the site since its creation time which is determined using DAS.

Having collected this information, $\langle N_{\text{replicas}} \rangle$ can be computed for a given time interval $[t_0, t_1]$. Then, summing over the sites:

$$\langle N_{\text{replicas}} \rangle = \sum_{S \in \text{sites}} \frac{\text{time on } S \text{ during } [t_0, t_1]}{t_1 - t_0} \quad (1)$$

This gives the average number of replicas of a dataset in a specific time interval. If a dataset was not at the site for the entire considered time interval the count is prorated with the fraction of time it was present in the interval.

2.3 Number of Accesses N_{accesses} and File N_{files} , and Size size

The remaining variables are relatively easily calculated. The number of files and size of a dataset are retrieved from DAS. It should be noted that we assume each replica is a "full" replica. An incomplete replica may be missing files and consequently it will have a smaller size. We compute N_{accesses} using the caches maintained by Detox. Detox is the Dynamic Data Management tool which deals with the deletion of deprecated datasets and under-utilized replicas. In order to make the latter decision, Detox keeps a local record of the number of accesses made to each replica of each dataset, derived from the popularity API. N_{accesses} for a dataset is defined as the total number of accesses over all replicas of that dataset.

3 Filling the plot

3.1 Dataset selection

We consider all datasets that are currently on, or have been on, Tier 2 sites (Tier-1s are excluded). This information is gathered from two places. First, we ask Phedex for all datasets which are currently on Tier-2s. Then, we use the `deleterrequests` Phedex API to determine all datasets which have been deleted during the relevant time interval. We consider all datasets in the union of these two sets. Finally, we double-check using the Phedex transfer/deletion histories that it actually was on at least one site during the interval. If the dataset passes this check, a corresponding entry is made in the histogram, as discussed below.

3.2 Binning

Having computed these variables for each dataset, the popularity plot may be made. The histogram is filled for each dataset by choosing the following bin-value:

$$\frac{N_{\text{accesses}}}{N_{\text{files}} \cdot \langle N_{\text{replicas}} \rangle} \quad (2)$$

The factor of N_{files} in the denominator is due to the fact that a single request to a dataset actually consists of a series of requests to each file in the dataset. Dividing by N_{files} ensures that this quantity is the same for small and large datasets. The entry is given weight:

$$\langle N_{\text{replicas}} \rangle \cdot \text{size} \quad (3)$$

For ease of comparing plots made under different conditions, the bin-value is normalized to the length of the time interval (in Figure 1, the unit of time is months). Finally, the plot is normalized to have an integral of unity. The un-normalized integral can be thought of as a measure of “average data volume” during the interval, since it can be computed as:

$$\sum_{\text{datasets}} \langle N_{\text{replicas}} \rangle \cdot \text{size} \quad (4)$$

Finally, it should be noted that there are two special bins in Figure 1. The very last bin is the overflow bin. The very first bin contains those entries for which $N_{\text{accesses}} = 0$ exactly. It is important to recall that the histogram is only filled with datasets that were on at least one site during the relevant interval. Thus, the very first bin shows the fractional volume of datasets which were on disk, but not accessed, during an interval.

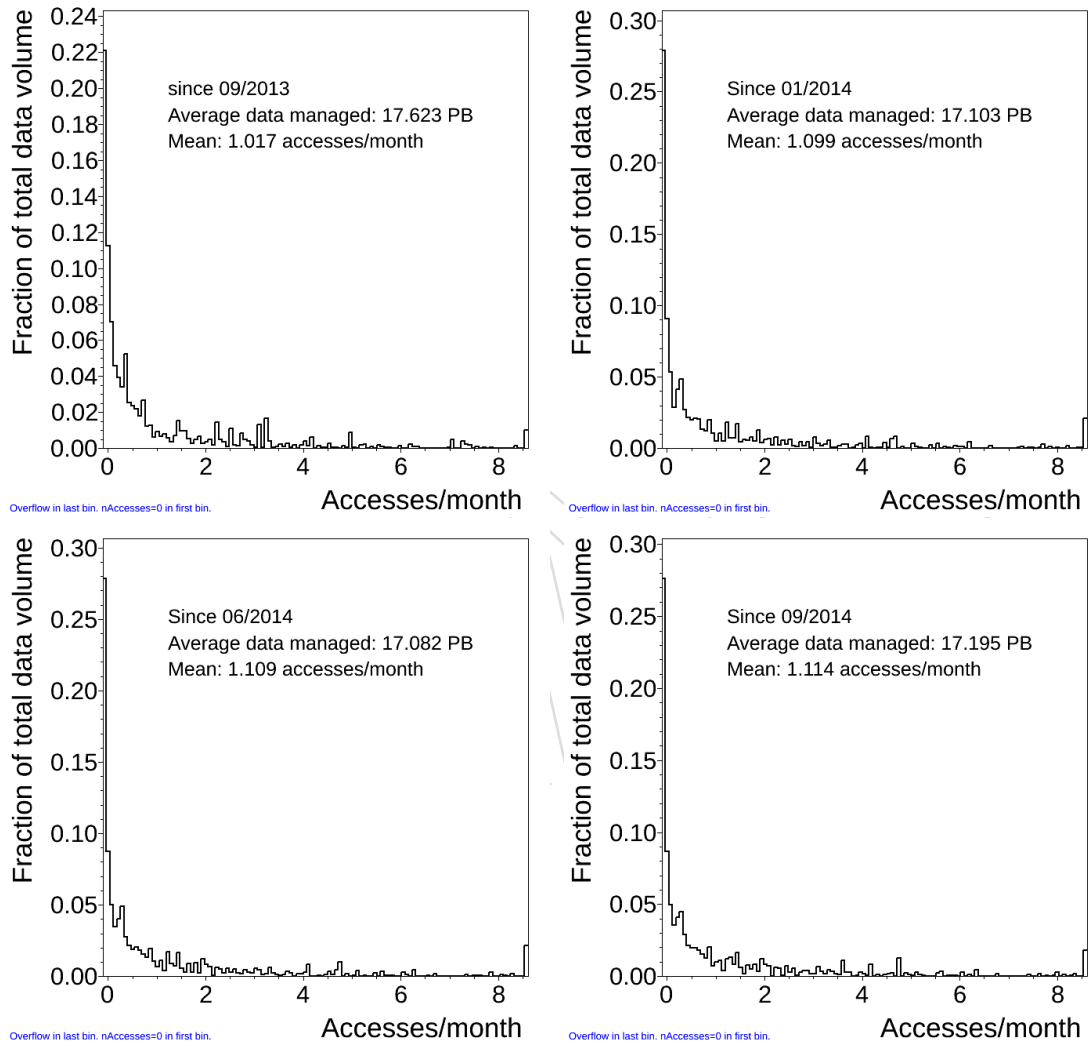


Figure 1: Usage plots datasets in AnalysisOps for various time intervals.

94 **References**

DRAFT